

Part 1: Theoretical Understanding

1. Short Answer Questions

Q1: Explain the primary differences between TensorFlow and PyTorch. When would you choose one over the other?

Answer:

TensorFlow and PyTorch are two of the most widely used deep learning frameworks today. While they serve similar purposes, they differ significantly in how they operate and are applied.

- **TensorFlow** (developed by Google) uses a **static computation graph**, meaning the model architecture is defined before execution. This approach offers performance benefits, especially in production environments. TensorFlow integrates well with:
 - **TensorBoard** (for visualization)
 - **TensorFlow Lite** (for mobile deployment)
 - **TensorFlow Serving** (for model deployment)
- **PyTorch** (developed by Facebook AI Research) uses a **dynamic computation graph**, which is built on-the-fly during execution. This makes PyTorch highly flexible and intuitive, particularly useful in research where models change frequently.

When to choose:

- ☒ Use **TensorFlow** for production environments, robust deployment, or mobile/edge devices.
- ☒ Use **PyTorch** for research, experimentation, and ease of debugging with Pythonic code.

Summary:

PyTorch is often preferred in **academia and research**, while TensorFlow is widely adopted in **industry and production** settings.

Q2: Describe two use cases for Jupyter Notebooks in AI development.

Answer:

Jupyter Notebooks are essential in AI and data science as they combine code, output, visualization, and documentation in a single document.

1. Exploratory Data Analysis (EDA):

Before training models, Jupyter allows step-by-step exploration of datasets using Python and visualization libraries like Matplotlib and Seaborn. This helps detect patterns, outliers, or class distributions, and allows documentation using markdown cells.

2. Rapid Prototyping and Experimentation:

Developers can test small code blocks incrementally, adjust hyperparameters, evaluate models, and visualize progress without leaving the notebook. This makes it ideal for quick iteration and collaborative sharing.

Q3: How does spaCy enhance NLP tasks compared to basic Python string operations?

Answer:

spaCy is a modern NLP library optimized for industrial use. Unlike basic string methods like `.split()`, `.find()`, or `.replace()`, spaCy brings linguistic intelligence to text processing.

Key Enhancements:

- **Context-Aware Tokenization:**

spaCy handles punctuation, contractions, and special entities correctly (e.g., treating "U.S.A." as one token).

- **Part-of-Speech Tagging and Dependency Parsing:**

spaCy identifies grammatical roles (noun, verb, adjective, etc.) and syntactic relationships between words.

- **Named Entity Recognition (NER):**

spaCy can extract key entities from text such as:

- "Apple" → *Organization*
- "Beats" → *Organization*
- "2014" → *Date*
- "\$3 billion" → *Money*

These features make spaCy ideal for chatbots, document classification, and information extraction — far beyond what simple string manipulation can achieve.

2. Comparative Analysis: Scikit-learn vs. TensorFlow

Aspect	Scikit-learn	TensorFlow
Target Applications	Best for classical ML tasks (e.g., regression, decision trees, SVM, clustering).	Best for deep learning tasks (e.g., CNNs, RNNs, large-scale neural networks).
Ease of Use	Very beginner-friendly with a consistent API (<code>.fit()</code> , <code>.predict()</code> , <code>.score()</code>).	Slightly complex; uses computational graphs. TensorFlow 2.x is more user-friendly with Keras.
Community Support	Mature and strong community; widely used in academia and traditional ML projects.	Very active community in deep learning, backed by Google with rich documentation and tools.

Summary:

- Use **Scikit-learn** for traditional ML and when starting out.
- Use **TensorFlow** for developing deep learning models, large-scale training, and production deployments.



Part 3: Ethics & Optimization

This section explores ethical considerations in AI models and demonstrates how to troubleshoot a faulty TensorFlow model.



1. Ethical Considerations



A. Potential Biases in Your Models

a) MNIST (Digit Classifier)

- **Bias Type:** Model bias due to **overrepresentation of certain digits** or **imbalanced data augmentation**.
- **Impact:** The model may perform better on digits like "1" and worse on ambiguous ones like "5" or "9".
- **Example:** A digit drawn sloppily or from a different cultural handwriting style (e.g., Arabic, left-handed) may be misclassified.

b) Amazon Reviews (NER + Sentiment)

- **Bias Type:** **Lexical bias** — based on limited positive/negative word lists.
- **Impact:** Fails to detect sarcasm, cultural nuance, or contextual meaning.
- **Example:** "This phone is *not bad at all*" would score as **negative**, though it's actually **positive**.

```
In [ ]:
```

✓ B. Mitigation Using AI Fairness Tools

✓ TensorFlow Fairness Indicators

- A tool to evaluate **fairness metrics** across different slices of data (e.g., age, gender, etc.)
- You could apply it in MNIST if metadata (e.g., country, school, or writing hand) were available.

```
In [3]: # Not required for this task, but here's an example usage idea:
# Evaluate performance disparity on digits drawn by left-handed vs. right-handed in
```

✓ spaCy Rule-Based Mitigation

- You can improve fairness and reduce sentiment misclassification by:
 - Expanding the sentiment word lists based on **domain-specific** vocabulary.
 - Using **dependency parsing** to detect negations.
 - Integrating with libraries like **TextBlob** or **VADER** for better handling of nuanced sentiment.

```
In [4]: # Example tweak to reduce bias in spaCy rule-based sentiment:
if "not good" in review.lower():
    neg += 1
    pos -= 1 # override naive match
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[4], line 2
      1 # Example tweak to reduce bias in spaCy rule-based sentiment:
----> 2 if "not good" in review.lower():
      3     neg += 1
      4     pos -= 1 # override naive match

NameError: name 'review' is not defined
```

✓ 2. Troubleshooting Challenge




Let's simulate a buggy TensorFlow script and show how to fix it.

✗ Buggy Code Snippet (Hypothetical)

```
In [ ]: model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(128, input_shape=(28, 28)))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5)
```

! What's Wrong?

Issue	Description
 Input Shape	<code>Dense</code> expects a 1D input, not a 2D image <code>(28, 28)</code>
 Loss Function	Using <code>categorical_crossentropy</code> but labels likely are not one-hot encoded
 Missing Flatten Layer	Need to flatten 2D input before Dense

✓ Fixed Version

```
In [ ]: model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)), # Fix input shape
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

# If using sparse labels (integers), use sparse_categorical_crossentropy
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['a

model.fit(x_train, y_train, epochs=5)
```